

# **Homework #2 Xv6**

Course

**CSE 460 Operating Systems**

Instructor

**Dr. Yan Zhang**

Meeting Time

**Mon. & Wed., 4:00 p.m. - 5:15 p.m.**

Due Date

**March 11, 2020**

Authors

**Kevin T. Vo**

**Esdras Lopez**

**Joseph Gonzales**

**Trevor Shortlidge**

**Brian Ayala**

## Group Members' Information

Name	Coyote ID	Responsibility
Kevin T. Vo	006316930	<ul style="list-style-type: none"><li>● Implemented date system call &amp; command (Part a) into xv6</li><li>● Implement part of uid, gid, ppid (Part c)</li><li>● Contributed to this report/documentation</li></ul>
Esdras Lopez	006198864	<ul style="list-style-type: none"><li>● Implemented part of uid, gid, ppid (Part c)</li><li>● Contributed to this report/documentation</li></ul>
Joseph Gonzales	006242648	<ul style="list-style-type: none"><li>● Implemented ps command and Ctrl+P (Part d)</li><li>● Contributed to this report/documentation</li></ul>
Trevor Shortlidge	006310209	<ul style="list-style-type: none"><li>● Implemented part of Ctrl+P (Part b)</li><li>● Reformatted output of ps command (Part d)</li><li>● Contributed to this report/documentation</li></ul>
Brian Ayala	006191688	<ul style="list-style-type: none"><li>● Implemented part of Ctrl+P (Part b)</li><li>● Contributed to this report/documentation</li></ul>

# Table of Contents

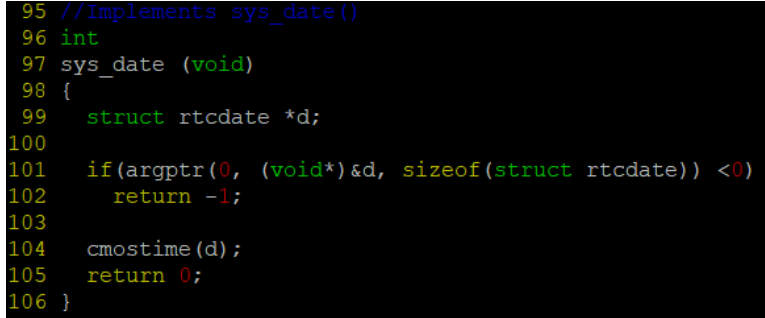
I.	Demonstration of the date() System Call .....	4
II.	Demonstration of the Ctrl-P Special Control Sequence.....	6
III.	Demonstration of UIDs, GIDs, and PPIDs.....	7
IV.	Demonstration of the “ps” Command.....	13

## I. Demonstration of the date() System Call

The system call “date” has been implemented to the Xv6 operating system where it would display the current day, month, year, hour, minute, and second in the format below:

*DAY/MONTH/YEAR HOUR:MINUTE:SECOND*

During the addition of this system call the following files have been created/modified:

Filename	Created / Modified	Line Numbers	Code
Makefile			
syscall.h	Modified	23 & 248	
user.h	Modified	26	
sysproc.c	Modified	95-106	
usys.S	Modified	32	
syscall.c	Modified	106 & 136	

date.c	Created	1-18	<pre> 1 #include "types.h" 2 #include "user.h" 3 #include "date.h" 4 5 int 6 main(int argc, char *argv[]) 7 { 8     struct rtcdate r; 9 10    if(date(&amp;r)) 11    { 12        printf(2, "date failed\n"); 13    } 14 15    printf(1, "%d/%d/%d %d:%d:%d\n", r.day, r.month, r.year, r.hour, r.minute, r.second); 16 17    exit(); 18 } </pre>
--------	---------	------	--

### Result from Execution of date() System Call:

```

SeaBIOS (version 1.12.0-2.fc30)

iPXE (http://ipxe.org) 00:03.0 C980 PCI2.10 PnP PMM+1FF91280+1FED1280 C980

Booting from Hard Disk...
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$
$
$ date
29/2/2020 0:5:14
$ █

```

Results from the syscall “date” comes from implementing the above into a full fledged syscall. Modified syscall.h that contains the mapping of the call number linked to SYS\_date. Modified defs.h to hold the kernel wide function call for date. Modified user.h that contains the functions prototype for date that is required to run user programs. Implemented the function for date inside of sysproc.c by getting the arguments off the stack. This implementation was provided to us. Next we modified the usys.S that exports the system call. Modified syscall.c to add the extern int for being in another file. Then the last step is to add the date.c file where we invoke the syscall implementation.

## II. Demonstration of the Ctrl-P Special Control Sequence

Result for Ctrl-P Special Control Sequence:

```
PID      Name   UID    GID    Elapsed    CPU    State    Size    PCs
1        init    0      0      278        1      sleep    12288   80103f57 80103ffd 80104cc9 80105e31 80105b7c
2        sh     0      0      276        0      sleep    16384   80103fic 801002da 8010101c 80104fb2 80104cc9 80105e31 80105b7c
```

Note: UID,GID, CPU are from part D, Dr.Zhang said it was ok to leave elapsed time in milliseconds.

Proc.h line 56 -> added uint start\_ticks to structure proc

```
uint start_ticks;
```

Modified procdump in proc.c to display our results in the ptable to output elapsed time, & size

Elapsed time result is done by subtracting ticks minus start\_ticks do give us the delta in milliseconds.

However, changing it to a float and dividing by 1000 to give us the result in seconds was not working with cprintf. Refer to to the comment I made before about Dr.Zhang letting our group use milliseconds instead.

Also p->sz is a built in variable in proc that deals with the size of each block of memory.

We just outputted the size : p->sz to display.

Proc.c line 131 -> initialized start\_ticks to ticks that is a global counter in milliseconds.

```
//added here the cp part for initializing start_ticks
p->start_ticks = ticks;
```


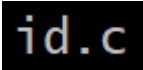
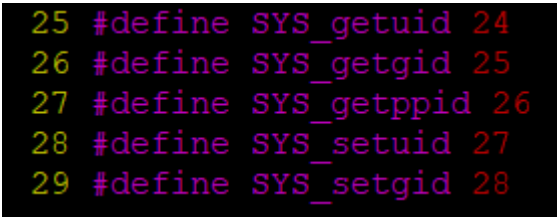
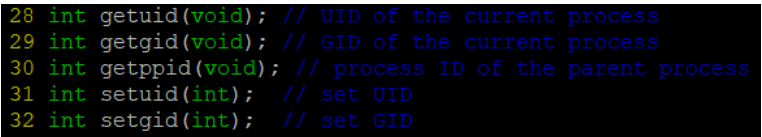
proc.c procdump() - line 540

```
printf("\n");
printf("PID \t Name \t UID \t GID \t Elapsed \t CPU \t State \t Size \t PCs \t\n");

for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    if(p->state == UNUSED)
        continue;
    if(p->state >= 0 && p->state < NELEM(states) && states[p->state])
        state = states[p->state];
    else
        state = "???";
    printf("%d \t %s \t %d \t %d \t %d \t %d \t %s \t %d",p->pid, p->name, p->uid, p->gid, (ticks - p->start_ticks), p->cpu_total, state, p->sz);
```

### III. Demonstration of UIDs, GIDs, and PPIDs

The implements the feature of users and groups into xv6 through uid and gid where it be used to store ID unsigned integers for tracking the ownership of a process for a user or group. By typing “id” into the terminal when executing xv6, the system call will invoke and demonstrate this feature by displaying the UID, GID and PPID that has been established by the test function within “id.c”.

Filename	Created / Modified	Line Numbers	Image
Makefile		178 & 248	 
syscall.h	Modified	25-29	
user.h	Modified	28-32	

proc.c	Modified	597-691	<pre> 597 //Set UID 598 int 599 setuid (uint uid) 600 { 601     //Checks if UID is in range to continue 602     if(uid &lt; 0    uid &gt; 32767) 603         return -1; 604     acquire(&amp;ptable.lock); 605     //const uint uid2 = 0; 606     struct proc *curproc = myproc(); 607     //if(argint(uid2, &amp;uid) &lt; 0) 608     // return -1; 609     if(uid &lt; 0 &amp;&amp; uid &gt; 32767) 610     { 611         //printf("\nError: Value for UID is &lt; 0 or &gt; 32,767\n"); 612         return -1; 613     } 614     curproc-&gt;uid = uid; 615     release(&amp;ptable.lock); 616     return 0; 617 } 618 619 //Set GID 620 int 621 setgid (uint gid) 622 { 623     //checks if gid is in range to continue 624     if(gid &lt; 0    gid &gt; 32767) 625         return -1; 626     acquire(&amp;ptable.lock); 627     struct proc *curproc = myproc(); 628     curproc-&gt;gid = gid; 629     release(&amp;ptable.lock); 630     return 0; 631 } 632 633 //Get UID of current process 634 uint 635 getuid(void) 636 { 637     struct proc *curproc = myproc(); 638     return curproc-&gt;uid; 639 } 640 641 //Get GID of current process 642 uint 643 getgid(void) 644 { 645     return myproc()-&gt;gid; 646 } 647 648 //Get process ID of parent 649 uint 650 getppid(void) 651 { 652     //struct proc *curproc = myproc(); 653     //uint ppid = myproc()-&gt;parent-&gt;pid; 654     //p-&gt;parent-&gt;pid; 655     struct proc *p; 656     sti(); 657     acquire(&amp;ptable.lock); 658     p = ptable.proc; 659     uint temp = p-&gt;pid; 660     release(&amp;ptable.lock); 661     if(temp &gt; 0) 662     { 663         return temp; 664     } 665     else 666     { 667         temp = 0; 668         return temp; 669     } 670 } </pre>
--------	----------	---------	--



sysproc.c	Modified	114-162	<pre> 114 //Returns the UID 115 int 116 sys_getuid(void) 117 { 118     return getuid(); 119 } 120 121 //Returns the GID 122 int 123 sys_getgid(void) 124 { 125     return getgid(); 126 } 127 128 //Returns the PPID 129 int 130 sys_getppid(void) 131 { 132     return getppid(); 133 } 134 135 136 //Set the UID with passing arguments into a kernel function 137 int sys_setuid(void) 138 { 139 140     int uid; 141 142     if(argint(0, &amp;uid) &lt; 0) 143         return -1; 144 145     //int setuid(uint) already checks if (uid &lt; 0    uid &gt; 32767) 146 147     return setuid(uid); 148 } 149 150 //Set GID with passing arguments into a kernel function 151 int sys_setgid(void) 152 { 153     int gid; 154 155     if(argint(0, &amp;gid) &lt; 0) 156         return -1; 157 158     //int setgid(uint) already checks for if (gid &lt; 0    gid &gt; 32767) 159     // return -1; 160 161     return setgid(gid); 162 } </pre>
usys.S	Modified	34-38	<pre> 34 SYSCALL(getuid) 35 SYSCALL(getgid) 36 SYSCALL(getppid) 37 SYSCALL(setuid) 38 SYSCALL(setgid) </pre>
syscall.c	Modified	108-112 & 138-142	<pre> 108 extern int sys_getuid(void); 109 extern int sys_getgid(void); 110 extern int sys_getppid(void); 111 extern int sys_setuid(void); 112 extern int sys_setgid(void); </pre>

			<pre> 138 [SYS_getuid]   sys_getuid, 139 [SYS_getgid]   sys_getgid, 140 [SYS_getppid]   sys_getppid, 141 [SYS_setuid]    sys_setuid, 142 [SYS_setgid]    sys_setgid, </pre>
id.c	Created	1-26	<pre> 1 #include "types.h" 2 #include "stat.h" 3 #include "user.h" 4 5 int 6 main(void) 7 { 8     uint uid, gid, ppid; 9 10    uid = getuid(); 11    printf(2, "Current UID is: %d\n", uid); 12    printf(2, "Setting UID to 100\n"); 13    setuid(100); 14    uid = getuid(); 15    printf(2, "Current UID is: %d\n", uid); 16    gid = getgid(); 17    printf(2, "Current GID is: %d\n", gid); 18    printf(2, "Setting GID to 100\n"); 19    setgid(100); 20    gid = getgid(); 21    printf(2, "Current GID is: %d\n", gid); 22    ppid = getppid(); 23    printf(2, "My parent process is: %d\n", ppid); 24    printf(2, "Done!\n"); 25    exit(); 26 } </pre>

Description of changes made within the file "syscall.h":

- The system calls responsible for setting and returning the UID, GID, and returning PPID are added so that it system call name can be mapped to a system call number so that it can be invoked within xv6.

Description of changes within the file "user.h":

- The prototype to the getter and setter functions for setting and returning the unsigned integers of UID and GID including one getter function for the PPID.
- The prototypes included are to the functions that holds the code to perform the task that a user uses a system call to request for.

Description of changes within the file “proc.c”:

- proc.c contains the functions:
  - uint getuid (void) // UID of the current process
  - uint getgid (void) // GID of the current process
  - uint getppid ( void) // process ID of the parent process
  - int setuid (uint) // set UID
  - int setgid (uint) // set GID

to actually sets the UID and GID using the setuid(uint) and setgid(uint) functions.

Note that user/xv6 will use these two functions to set the UID and GID. The getter functions uint getuid(void), uint getgid(void), and uint getppid(void) that returns the UID, GID, and PPID respectively. It is within setter functions where it checks if the value UID and GID are within range of [0, 32727].

Description of changes within the file “sysproc.c”:

- Included the sys functions that returns the actual functions listed above to perform the process associated its system call.

Description of changes within the file “usys.S”:

- “usys.S” contains the system calls to be made available by the kernel:

```
34 SYSCALL(getuid)
35 SYSCALL(getgid)
36 SYSCALL(getppid)
37 SYSCALL(setuid)
38 SYSCALL(setgid)
```

Description of changes within the file “syscall.c”:

- First contains the system call entry points:

```
108 extern int sys_getuid(void);
109 extern int sys_getgid(void);
110 extern int sys_getppid(void);
111 extern int sys_setuid(void);
112 extern int sys_setgid(void);
```

- Secondly, “syscall.c” contains a dispatch table that defines from the symbol name made earlier within files “usys.S” and “syscall.h” to its function name. All of this is done within the “int (\*syscalls[])(void)” within the “syscall.c”.

Description of changes made within file “id.c”:

- Contains the test function to run the system calls implemented in this part.

Result from UIDs, GIDs, and PPIDs:

```
SeaBIOS (version 1.12.0-2.fc30)

iPXE (http://ipxe.org) 00:03.0 C980 PCI2.10 PnP PMM+1FF91280+1FED1280 C980

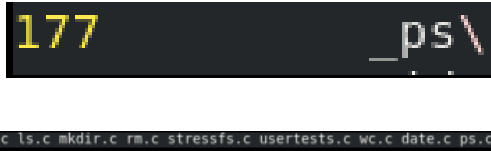

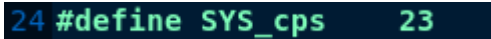
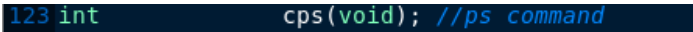
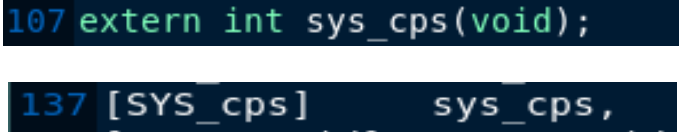
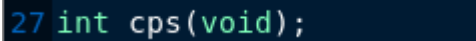
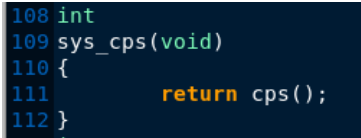

Booting from Hard Disk...
xv6...
cpul: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ id
Current UID is: 0
Setting UID to 100
Current UID is: 100
Current GID is: 0
Setting GID to 100
Current GID is: 100
My parent process is: 1
Done!
$ █
```

The following results derived from the test function within “id.c” for testing the `uint getuid(void)`, `uint getgid(void)`, `uint getppid(void)`, `int setuid(uint)`, and `int setgid(uint)` functions along with the system call to them. The `setuid(uint)` and `setgid(uint)` takes in an unsigned integer so assigns them the UID or GID for a process. The getters `uint getuid(void)` and `uint getgid(void)` returns value of the UID and GID that has been set. Lastly, `uint getppid(void)` returns back the parent process by first calling the `sti()` to interrupt the processor. Then, it acquires the `p` table which contains all the process. However, this function will only focus on the parent process with `uint temp = p->pid`; and then returns it. If the `p->pid` happens to be a negative value, then `temp = 0`.

## IV. Demonstration of the “ps” Command

The system call “cps” has been implemented to the Xv6 operating system in order to display the ptable for the current processes, through the use of the “ps” command. It outputs the following process information: pid, name, UID, GID, PPID, ELAPSED time (in ms), CPU time (in ms), size, and state.

During the addition of this system call the following files have been created/modified:

Filename	Created / Modified	Line Numbers	Image
Makefile	Modified	177 & 248	
ps.c	Created	1-12	
syscall.h	Modified	24	
defs.h	Modified	123	
syscall.c	Modified	107 & 137	
user.h	Modified	27	
sysproc.c	Modified	108-112	
usys.S	Modified	33	

proc.c	Modified	707-740	<pre> 707 //current process status 708 int 709 cps() 710 { 711     struct proc *p; 712 713     // Enable interrupts on this processor. 714     sti(); 715 716     // Loop over process table looking for process with pid. 717     acquire(&amp;ptable.lock); 718     cprintf("pid %d name %s uid %d gid %d elapsed %d CPU %d size %d state\n", 719             p-&gt;pid, p-&gt;name, p-&gt;uid, p-&gt;gid, p-&gt;elapsed, p-&gt;cpu, p-&gt;size, p-&gt;state); 720     for(p = ptable.proc; p &lt; &amp;ptable.proc[NPROC]; p++) 721     { 722         if(p-&gt;parent-&gt;pid &lt; 0) 723             p-&gt;parent-&gt;pid = 0; 724         if(p-&gt;state == SLEEPING) 725             cprintf("pid %d name %s uid %d gid %d elapsed %d CPU %d size %d state\n", 726                     p-&gt;pid, p-&gt;name, p-&gt;uid, p-&gt;gid, p-&gt;elapsed, p-&gt;cpu, p-&gt;size, p-&gt;state); 727         else if(p-&gt;state == RUNNING) 728             cprintf("pid %d name %s uid %d gid %d elapsed %d CPU %d size %d state\n", 729                     p-&gt;pid, p-&gt;name, p-&gt;uid, p-&gt;gid, p-&gt;elapsed, p-&gt;cpu, p-&gt;size, p-&gt;state); 730         else if(p-&gt;state == UNRUNNABLE) 731             cprintf("pid %d name %s uid %d gid %d elapsed %d CPU %d size %d state\n", 732                     p-&gt;pid, p-&gt;name, p-&gt;uid, p-&gt;gid, p-&gt;elapsed, p-&gt;cpu, p-&gt;size, p-&gt;state); 733         else if(p-&gt;state == ZOMBIE) 734             cprintf("pid %d name %s uid %d gid %d elapsed %d CPU %d size %d state\n", 735                     p-&gt;pid, p-&gt;name, p-&gt;uid, p-&gt;gid, p-&gt;elapsed, p-&gt;cpu, p-&gt;size, p-&gt;state); 736     } 737     release(&amp;ptable.lock); 738     return 0; 739 } </pre>
--------	----------	---------	--

The total CPU time was completed as follows:

proc.h lines 56 & 57 -> added uint cpu\_total and uint cpu\_runtime to the proc structure.

```

56     uint cpu_total;           //total amount of ticks the cpu had done
57     uint cpu_runtime;        //runtime of cpu

```

proc.c line 132 -> cpu\_total initialized to 0.

```

132     p->cpu_total = 0;

```

proc.c line 376 -> cpu\_runtime set to ticks.

```

375     //adding up the cpu process time
376     p->cpu_runtime = ticks;

```

proc.c line 413 -> cpu\_total calculated by adding (ticks minus the cpu\_runtime) to give us the total CPU time in milliseconds.

```

412     //getting the correct ammount of time(delta) for the process to finish!
413     p->cpu_total += ( ticks - p->cpu_runtime );

```

### Result from Execution of cps() System Call:

ps								
pid	name	UID	GID	PPID	ELAPSED	CPU	SIZE	STATE
1	init	0	0	0	283	36	12288	SLEEPING
2	sh	0	0	1	242	3	16384	SLEEPING
3	ps	0	0	2	5	1	12288	RUNNING

Results from the system call “cps” comes from implementing it, as well as the ps command, into Xv6. The “ps” command was created by creating the ps.c file and modifying the Makefile, in order to add the command to Xv6. Modification of syscall.h was done, mapping the call number linked to SYS\_cps. Additionally, defs.h was modified to hold the function call for cps. Next, syscall.c was modified in order to include the extern int of sys\_cps. A function prototype for cps was also added to user.h, used by user programs. Next, sysproc.c was modified to implement the function for cps. Exporting the system call was done by modifying usys.S to include the cps system call. Finally, the implementation for cps was added to proc.c, allowing the various process information to be output in a ptable.